

Lab 6 - Create a Service Spanning Multiple Containers

Set-up:

For the set up of this lab, I followed the steps given in the lab 6 document.

1. Create user-defined Docker network (named pgnet)

\$ docker network create pgnet

```
(env) dix@dix-ubuntu:cloud-sec-2019-06$ docker network create pgnet
d92fe406dd2101605e98faabcec2e19100ae0c810d822d8ea5ffc345126a6574
(env) dix@dix-ubuntu:cloud-sec-2019-06$
```

2. Postgres Container

- a. Pull the postgres image from Docker repositories

\$ docker pull postgres

```
(env) dix@dix-ubuntu:cloud-sec-2019-06$ docker pull postgres
Using default tag: latest
latest: Pulling from library/postgres
Digest: sha256:a4a944788084a92bcaff6180833428f17cceb610e43c828b3a42345b33a608a7
Status: Image is up to date for postgres:latest
docker.io/library/postgres:latest
(env) dix@dix-ubuntu:cloud-sec-2019-06$
```

- b. Run container using postgres image. Connect it to your Docker network

*\$ docker run --network pgnet --name mypg *

> -e POSTGRES_PASSWORD=password -d postgres

```
(env) dix@dix-ubuntu:cloud-sec-2019-06$ docker run \
> --network pgnet \
> --name mypg \
> -e POSTGRES_PASSWORD=password \
> -d \
> postgres
333d8b83d2cb8a8936ebff3b95bc5f5d7959ced27d7aefc1aa00f1e990809e64
(env) dix@dix-ubuntu:cloud-sec-2019-06$
```

3. Init Container

- a. Find IP address for postgres database container

*\$ docker container inspect mypg *

> -f '{{.NetworkSettings.Networks.pgnet.IPAddress}}'

```
(env) dix@dix-ubuntu:cloud-sec-2019-06$ docker container inspect mypg \
> -f '{{.NetworkSettings.Networks.pgnet.IPAddress}}'
172.18.0.2
(env) dix@dix-ubuntu:cloud-sec-2019-06$
```

- b. Write *init.sql* to build table in database

```
(env) dix@dix-ubuntu:cloud-sec-2019-06$ cat init.sql
CREATE TABLE IF NOT EXISTS pathcount (
  path TEXT PRIMARY KEY,
  count INT DEFAULT 0
);
(env) dix@dix-ubuntu:cloud-sec-2019-06$
```

- c. Run init container to initialize your database for use

```
$ docker run -i --rm --network pgnet -e PGPASSWORD=password postgres \
> psql -h 172.18.0.2 -U postgres < init.sql
```

```
(env) dix@dix-ubuntu:cloud-sec-2019-06$ docker run \
> -i \
> --rm \
> --network pgnet \
> -e PGPASSWORD=password \
> postgres \
> psql -h 172.18.0.2 -U postgres < init.sql
CREATE TABLE
(env) dix@dix-ubuntu:cloud-sec-2019-06$
```

- d. I did run into one error when following the lab document. When trying to run the init container in interactive mode using `--tty` (or `-it`), I would get an 'Input device is not a TTY' error. The only work around I could find without digging too deep was to just omit `--tty` flag. Screenshot of error below:

```
(env) dix@dix-ubuntu:cloud-sec-2019-06$ docker run \
> -it \
> --rm \
> --network pgnet \
> -e PGPASSWORD=password \
> postgres \
> psql -h 172.18.0.2 -U postgres < init.sql
the input device is not a TTY
(env) dix@dix-ubuntu:cloud-sec-2019-06$
```

Service Container:

1. For my service container, the package requirements include the following python modules; installed using `$ pip install [package]`:
 1. Flask for creating my web application
 2. Python-dotenv for handling virtual environment variables
 3. Psycopg2 to work with postgresql databases.

Dependencies for psycopg2 included the postgresql package. The current version in the Ubuntu 18.04 apt repositories is 10. With `$ sudo apt-get install`, I installed the two dependencies:

1. postgresql
 2. postgresql-server-dev-10
2. I wrote my web service application using Python and flask. I have all paths going to the default route which calls functions that will query and update the postgres database and display all paths onto an HTML page.
 - a. As required, I defined all configuration variables in ENVIRONMENT VARIABLES. To use python-dotenv, I defined those variables in a '.env' file in the project directory. That file is then loaded in 'main.py'. Here is that file:

```
(env) dix@dix-ubuntu:cloud-sec-2019-06$ cat .env
# Environment variables for Docker lab 06
POSTGRES_USER="postgres"
POSTGRES_DATABASE="postgres"
POSTGRES_PASSWORD="password"
POSTGRES_HOST="172.18.0.2"
(env) dix@dix-ubuntu:cloud-sec-2019-06$
```

- b. The primary code files include:
 - i. [Main.py](#) : web application file
 - ii. [Index.html](#) : HTML for displaying path counts
 - iii. [Style.css](#) : styling for the web page

All of which are located in my class github:

https://github.com/richardbradt/cloud_computing/tree/master/cloud-sec-2019-06

3. I then built my container using the following command and Dockerfile:
`$ docker build -t pathcount .`

```

1 FROM python
2 RUN mkdir /app
3 COPY . /app/
4 WORKDIR /app
5 RUN pip install Flask \
6     python-dotenv \
7     psycpg2
8 EXPOSE 8080
9 CMD ["python", "main.py"]

```

Dockerfile for Service Container

- I then ran my service container with the following command, exposing port 8080 to host port 9090.

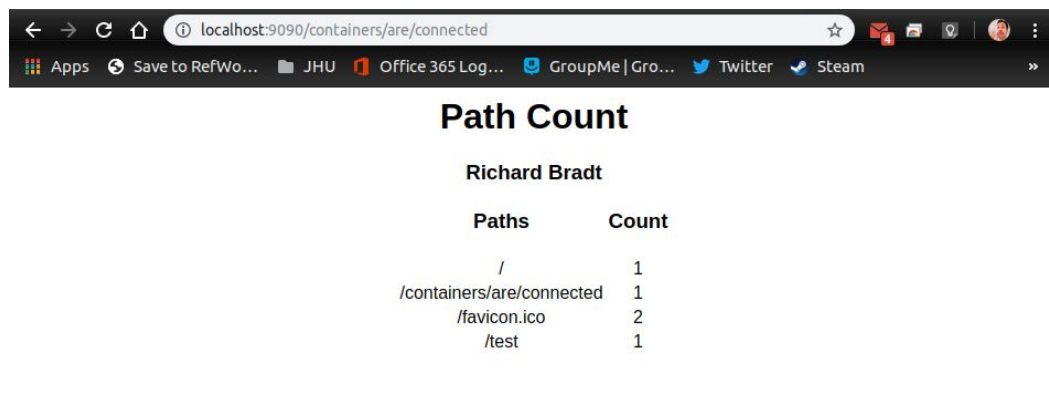
```
$ docker run -it --name=pathcount_cont --network pgnet \
> -p 9090:8080 pathcount
```

```

(env) dix@dix-ubuntu:cloud-sec-2019-06$ docker run \
> -it \
> --name=pathcount_cont \
> --network pgnet \
> -p 9090:8080 \
> pathcount
* Serving Flask app "main" (lazy loading)
* Environment: production
  WARNING: This is a development server. Do not use it in a production deployment.
  Use a production WSGI server instead.
* Debug mode: off
* Running on http://0.0.0.0:8080/ (Press CTRL+C to quit)

```

- Finally, going to <http://localhost:9090> on my host machine's browser produces this page, and by manually typing various paths in the address bar, the service continues to update and display the database contents.



The screenshot shows a web browser window with the address bar set to `localhost:9090/containers/are/connected`. The page displays the title "Path Count" and the name "Richard Bradt". Below this is a table with two columns: "Paths" and "Count". The table contains the following data:

Paths	Count
/	1
/containers/are/connected	1
/favicon.ico	2
/test	1

Questions:

1. Why did we create a special network instead of exposing the host network?
This further isolates the containers from elements on the host network.
2. Why didn't we use exposed ports everywhere (that they exist)?
By not exposing ports on the postgres container, we're isolating the database to our user-defined docker bridge network, protecting it from outside influences.
3. What could happen if you didn't use SQL parameters, but relied on string formatting for setting the path in your queries?
We would be subject to SQL Injection attacks.
4. Why is that particularly important in this setup? What makes those parameters potentially dangerous?
Parameters could still be dangerous if they're not sanitized.
5. The bridge network we define only works on a single host. What would you have to do to make these containers talk to each other if they were running on *different host machines*?
Docker would have to use the host network as opposed to using a user-defined bridge network like we did in this lab.
6. What parts of this did you wish were simpler? Which parts seemed unnecessarily difficult?
The most difficult portions of this lab falls to whichever part is most unfamiliar. For me, it was wrapping my head around the 'docker-isms'; most notably, understanding docker networks and docker postgres. Finding usable python modules for postgresql was just a Google search. Luckily, I found very useful articles and tutorials. But, sifting through docker documentation can be daunting and somewhat lacking.

Main.py

```
1  """App that produces a path count."""
2  """Default route will increment count for each path and store data in database."""
3
4  from flask import Flask, request, render_template
5  from os.path import join, dirname
6  from dotenv import load_dotenv
7  import os, psycopg2
8
9  """Import Environment Variables for DB"""
10 dotenv_path = join(dirname(__file__), '.env')
11 load_dotenv(dotenv_path)
12
13 user = os.getenv('POSTGRES_USER')
14 db = os.getenv('POSTGRES_DATABASE')
15 secret = os.getenv('POSTGRES_PASSWORD')
16 host = os.getenv('POSTGRES_HOST')
17
18 app = Flask(__name__)
19
20 """Default route will get path and check for database entry"""
21 """If in database, increment count. If not, set count to 1 and add to db"""
22 @app.route('/', defaults={'path': '/'}, methods=['GET'])
23 @app.route('/<path:path>', methods=['GET'])
24 def root(path):
25     c_path = request.path
26     print('CURRENT PATH: {}'.format(c_path))
27     count_path(c_path)
28     return display_paths()
29
30 """Check database for path and count"""
31 """Connect to Postgresql database and query current path"""
32 def count_path(path):
33     sql = """INSERT INTO pathcount (path, count)
34             VALUES (%s, 1)
35             ON CONFLICT (path) DO UPDATE
36             SET count = pathcount.count + 1
37             RETURNING count;"""
38     conn = None
39
40     try:
41         conn = psycopg2.connect(host=host, database=db, user=user, password=secret)
42         cur = conn.cursor()
43         cur.execute(sql, (path,))
44         conn.commit()
45         cur.close()
46     except (Exception, psycopg2.DatabaseError) as error:
47         print(error)
48     finally:
49         if conn is not None:
50             conn.close()
```



```

51
52 """Display current database contents"""
53 """Query Postgresql database and build JSON of all pathcounts"""
54 def display_paths():
55     sql = """SELECT path, count FROM pathcount ORDER BY path"""
56     conn = None
57
58     try:
59         conn = psycopg2.connect(host=host ,database=db, user=user, password=secret)
60         cur = conn.cursor()
61         cur.execute(sql)
62         entities = cur.fetchall()
63         path_json = build_json(entities)
64         conn.commit()
65         cur.close()
66     except (Exception, psycopg2.DatabaseError) as error:
67         print(error)
68     finally:
69         if conn is not None:
70             conn.close()
71
72     return render_template('index.html', data=path_json)
73
74 """Builds JSON to send to HTML/JS. Takes entity list from DB"""
75 def build_json(path_list):
76     new_json = '['
77     counter = 1
78     for ent in path_list:
79         print("CURRENT KEY: {}".format(ent[0]))
80         print("CURRENT KEY VALUE: {}".format(ent[1]))
81         if counter==len(path_list):
82             new_json+='{"path":"%s","count": "%d"}]' % (ent[0], ent[1])
83             break
84         else:
85             new_json+='{"path":"%s","count": "%d"},' % (ent[0], ent[1])
86             counter+=1
87
88     return new_json
89
90 if __name__ == '__main__':
91     app.run(host='0.0.0.0', port='8080')

```

Index.html

```
1 |
2 <html>
3 <head>
4 <title>Richard Bradt - Cloud Sec 2019</title>
5 <link type="text/css" rel="stylesheet" href="{{ url_for('static', filename='style.css') }}"></link>
6 </head>
7 <body>
8 <div id="main" class="mainDiv">
9   <h1>Path Count</h1>
10   <h3>Richard Bradt</h3>
11   <div class="container" id="table_div">
12     <table class="container" id="path_table">
13       <script>
14         var paths = JSON.parse('{{ data | safe }}');
15         let html = '<tr><th><h3>Paths</h3></th><th><h3>Count</h3></th></tr>';
16         for (let x of paths) {
17           html += '<tr><td>' + x.path + '</td><td>' + x.count + '</td></tr>';
18         }
19         document.getElementById('path_table').innerHTML = html;
20       </script>
21     </table>
22   </div>
23 </div>
24 </body>
25 </html>
26
```

Style.css

```
1 .mainDiv {
2   margin-left: auto;
3   margin-right: auto;
4   max-width: 1000px;
5   float: none;
6   font-family: "helvetica", sans-serif;
7   text-align: center;
8 }
9 .container {
10   text-align: center;
11 }
12 .container table {
13   margin: 0 auto;
14 }
15
```